

```
In [1]: import pandas as pd
```

```
In [1]: import numpy as np
```

```
In [3]: import matplotlib as mpl
```

```
In [4]: import matplotlib.pyplot as plt
```

```
In [5]: import seaborn as sns
```

```
In [4]: #1.1
import pandas as pd
weights_series = pd.read_csv(r"C:\Users\gabed\.ipython\weights.tsv")
print(weights_series)
```

```
      164
0    158
1    172
2    153
3    144
4    156
5    189
6    163
7    134
8    159
9    143
10   176
11   177
12   162
13   141
14   151
15   182
16   185
17   171
18   152
```

```
In [8]: #1.1
data = pd.Series([164, 158,
172, 153, 144, 156, 189, 163, 134, 159, 143, 176, 177, 162, 141, 151, 182, 185, 171,
152])
data
```

```
Out[8]: 0      164
1      158
2      172
3      153
4      144
5      156
6      189
7      163
8      134
9      159
10     143
11     176
12     177
13     162
14     141
15     151
16     182
17     185
18     171
19     152
dtype: int64
```

```
In [2]: #1.2
import pandas as pd

weight_lbs = ([164, 158,
172, 153, 144, 156, 189, 163, 134, 159, 143, 176, 177, 162, 141, 151, 182, 185, 171,
152])

weight_kg = [round(weight * 0.453592,2) for weight in weight_lbs]

weight_series = pd.Series(weight_kg)

print(weight_series)
```

```
0    74.39
1    71.67
2    78.02
3    69.40
4    65.32
5    70.76
6    85.73
7    73.94
8    60.78
9    72.12
10   64.86
11   79.83
12   80.29
13   73.48
14   63.96
15   68.49
16   82.55
17   83.91
18   77.56
19   68.95
dtype: float64
```

```
In [39]: #1.3
import pandas as pd

weight_lbs = ([164, 158,
172, 153, 144, 156, 189, 163, 134, 159, 143, 176, 177, 162, 141, 151, 182, 185, 171,
152])

weight_series_lbs = pd.Series(weight_lbs)

weight_kg = weight_series_lbs * 0.453592

mean_lbs = weight_series_lbs.mean()
median_lbs = weight_series_lbs.median()
std_dev_lbs = weight_series_lbs.std()

mean_kg = weight_kg.mean()
median_kg = weight_kg.median()
std_dev_kg = weight_kg.std()

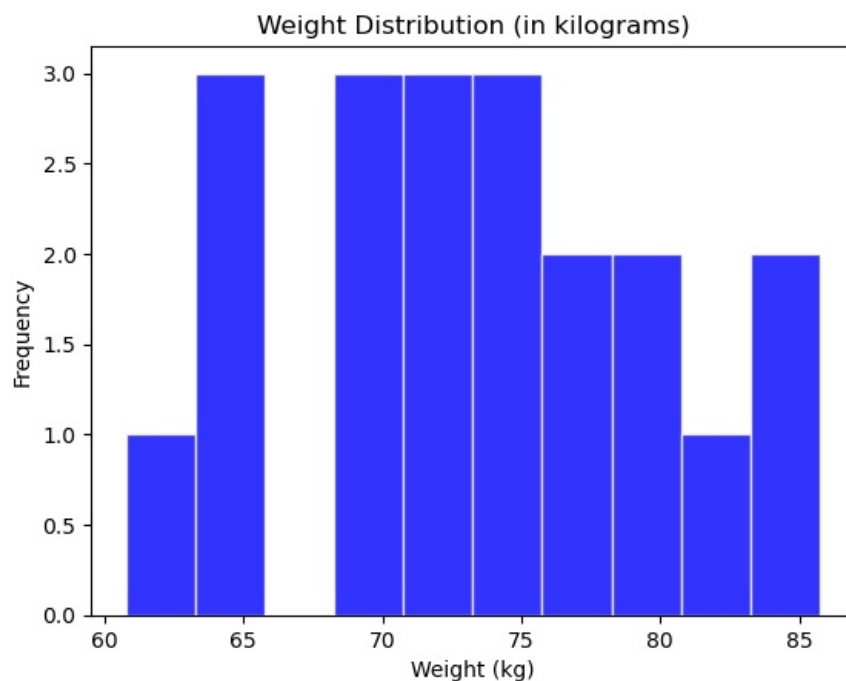
print("weight in lbs")
print("Mean:", mean_lbs)
print("Median:", median_lbs)
print("Standard Deviation:", std_dev_lbs)
print("weight in kg:")
print("Mean:", mean_kg)
print("Median:", median_kg)
print("Standard Deviation:", std_dev_kg)

weight in lbs
Mean: 161.6
Median: 160.5
Standard Deviation: 15.44906742203316
weight in kg:
Mean: 73.30046720000001
Median: 72.80151599999999
Standard Deviation: 7.007573390094864
```

```
In [20]: #1.4
import matplotlib.pyplot as plt

weights_lbs = [164, 158, 172, 153, 144, 156, 189, 163, 134, 159, 143, 176, 177, 162, 141, 151, 182, 185, 171, 1
weights_kg = [weight * 0.453592 for weight in weights_lbs]

plt.hist(weights_kg, bins=10, color='blue', edgecolor='white', alpha=0.8)
plt.title('Weight Distribution (in kilograms)')
plt.xlabel('Weight (kg)')
plt.ylabel('Frequency')
plt.show()
```




```
In [18]: #2.1
import pandas as pd
df = pd.read_csv(r"C:\Users\gated\.ipython\boston.csv")
num_rows, num_cols = df.shape
print("Number of rows:", num_rows)
print("Number of columns:", num_cols)
```

Number of rows: 506
Number of columns: 13

```
In [14]: #2.2
import pandas as pd
df = pd.read_csv(r"C:\Users\gated\.ipython\boston.csv")
index_lowest_nox = df['NOX'].idxmin()
medv_lowest_nox = df.loc[index_lowest_nox, 'MEDV']
print("owner-occupied home value (MEDV) for the lowest nitric oxide concentration (NOX):", medv_lowest_nox)
```

owner-occupied home value (MEDV) for the lowest nitric oxide concentration (NOX): 20.1

```
In [19]: #2.3
import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv(r"C:\Users\gated\.ipython\boston.csv")
plt.figure(figsize=(8, 8))
plt.boxplot(df['CRIM'], vert=False)
plt.title('Boxplot of Per Capita Crime Rate (CRIM)')
plt.xlabel('CRIM')
plt.show()
Q1 = df['CRIM'].quantile(0.25)
Q3 = df['CRIM'].quantile(0.75)
IQR = Q3 - Q1
print("Interquartile Range (IQR) for Crime Rate (CRIM):", IQR)
```



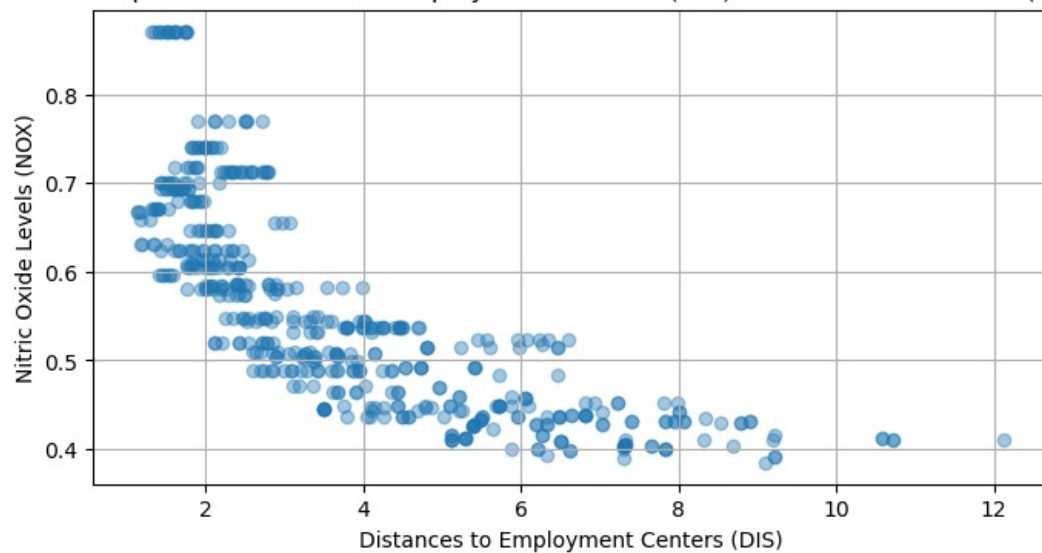
Interquartile Range (IQR) for Crime Rate (CRIM): 3.5950375

```
In [11]: #2.4
import pandas as pd
df = pd.read_csv(r"C:\Users\gabed\.ipython\boston.csv")
Q1 = df['CRIM'].quantile(0.25)
Q3 = df['CRIM'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
outliers_df = df[(df['CRIM'] < lower_bound) | (df['CRIM'] > upper_bound)]
mean_age_original = df['AGE'].mean()
mean_age_outliers = outliers_df['AGE'].mean()
print("Mean AGE in original dataframe:", mean_age_original)
print("Mean AGE in subsetted dataframe (with outliers):", mean_age_outliers)
```

Mean AGE in original dataframe: 68.57490118577078
Mean AGE in subsetted dataframe (with outliers): 94.23333333333335

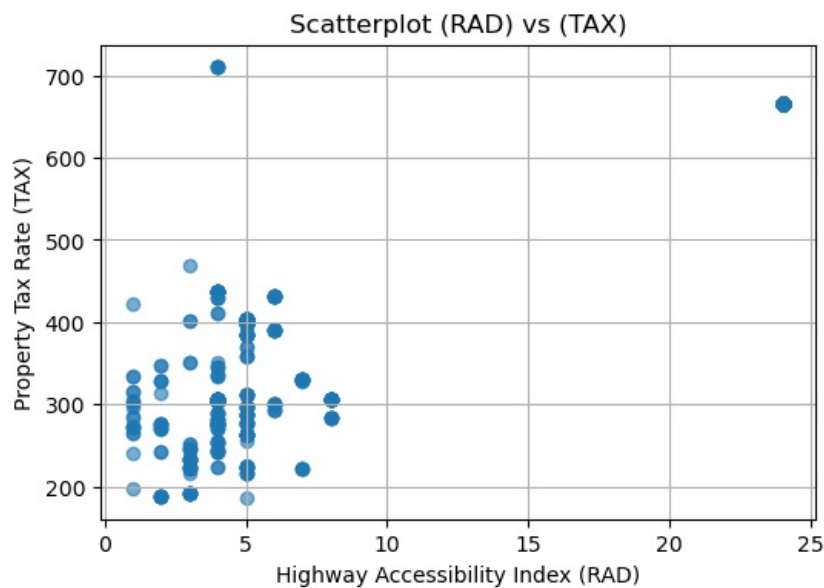
```
In [13]: #2.5
import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv(r"C:\Users\gabed\.ipython\boston.csv")
plt.figure(figsize=(8, 4))
plt.scatter(df['DIS'], df['NOX'], alpha=0.4)
plt.title('Scatterplot of Distances to Employment Centers (DIS) vs Nitric Oxide Levels (NOX)')
plt.xlabel('Distances to Employment Centers (DIS)')
plt.ylabel('Nitric Oxide Levels (NOX)')
plt.grid(True)
plt.show()
correlation_index = df['DIS'].corr(df['NOX'])
print("Correlation index between DIS and NOX:", correlation_index)
```

Scatterplot of Distances to Employment Centers (DIS) vs Nitric Oxide Levels (NOX)




Correlation index between DIS and NOX: -0.7692301132258278

```
In [5]: #2.6
import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv(r"C:\Users\gated\.ipython\boston.csv")
plt.figure(figsize=(6, 4))
plt.scatter(df['RAD'], df['TAX'], alpha=0.6)
plt.title('Scatterplot (RAD) vs (TAX)')
plt.xlabel('Highway Accessibility Index (RAD)')
plt.ylabel('Property Tax Rate (TAX)')
plt.grid(True)
plt.show()
correlation_index = df['RAD'].corr(df['TAX'])
print("Correlation index between RAD and TAX:", correlation_index)
upper_bound = df['TAX'].quantile(0.75) + 1.5 * (df['TAX'].quantile(0.75) - df['TAX'].quantile(0.25))
outlier_condition = df['TAX'] > upper_bound
df_cleaned = df[~outlier_condition]
plt.figure(figsize=(6, 4))
plt.scatter(df_cleaned['RAD'], df_cleaned['TAX'], alpha=0.6)
plt.title('Scatterplot (RAD) vs (TAX) without Outliers')
plt.xlabel('Highway Accessibility Index (RAD)')
plt.ylabel('Property Tax Rate (TAX)')
plt.grid(True)
plt.show()
correlation_index_cleaned = df_cleaned['RAD'].corr(df_cleaned['TAX'])
print("Correlation index between RAD and TAX after removing outliers:", correlation_index_cleaned)
```



Correlation index between RAD and TAX: 0.9102281885331835



Correlation index between RAD and TAX after removing outliers: 0.9102281885331835

```
In [17]: #3.1
import seaborn as sns
tips_df = sns.load_dataset('tips')
tips_df.head()
tips_df['tip_percent'] = (tips_df['tip'] / tips_df['total_bill']) * 100
tips_df['tip_percent'] = tips_df['tip_percent'].round(2)
print(tips_df.head())
```

	total_bill	tip	sex	smoker	day	time	size	tip_percent
0	16.99	1.01	Female	No	Sun	Dinner	2	5.94
1	10.34	1.66	Male	No	Sun	Dinner	3	16.05
2	21.01	3.50	Male	No	Sun	Dinner	3	16.66
3	23.68	3.31	Male	No	Sun	Dinner	2	13.98
4	24.59	3.61	Female	No	Sun	Dinner	4	14.68

```
In [8]: #3.2
import seaborn as sns
tips_df = sns.load_dataset('tips')
mean_bill_per_day = tips_df.groupby('day')['total_bill'].mean()
day_highest_mean_bill = mean_bill_per_day.idxmax()
highest_mean_bill = mean_bill_per_day.max()
print("Days in the dataset:", mean_bill_per_day.index.tolist())
print("Day with the highest bill mean:", day_highest_mean_bill)
print("Highest mean bill amount:", highest_mean_bill)
```

Days in the dataset: ['Thur', 'Fri', 'Sat', 'Sun']
Day with the highest bill mean: Sun
Highest mean bill amount: 21.41


```
In [9]: #3.3
import seaborn as sns
import pandas as pd
tips_df = sns.load_dataset('tips')
time_counts = tips_df['time'].value_counts()
smokers_during_lunch = tips_df[tips_df['time'] == 'Lunch']['smoker'].value_counts()
smokers_during_dinner = tips_df[tips_df['time'] == 'Dinner']['smoker'].value_counts()
ime_of_day_df = pd.DataFrame({'time_of_day': time_counts.index, 'count': time_counts.values})
smokers_lunch_df = pd.DataFrame({'smoker': smokers_during_lunch.index, 'count': smokers_during_lunch.values})
smokers_dinner_df = pd.DataFrame({'smoker': smokers_during_dinner.index, 'count': smokers_during_dinner.values})
merged_df = pd.merge(smokers_lunch_df, smokers_dinner_df, on='smoker', suffixes=('_lunch', '_dinner'))
merged_df['percent_lunch'] = (merged_df['count_lunch'] / time_counts['Lunch']) * 100
merged_df['percent_dinner'] = (merged_df['count_dinner'] / time_counts['Dinner']) * 100
print("Percentage of smokers during lunch and dinner:")
print(merged_df)
```

Percentage of smokers during lunch and dinner:

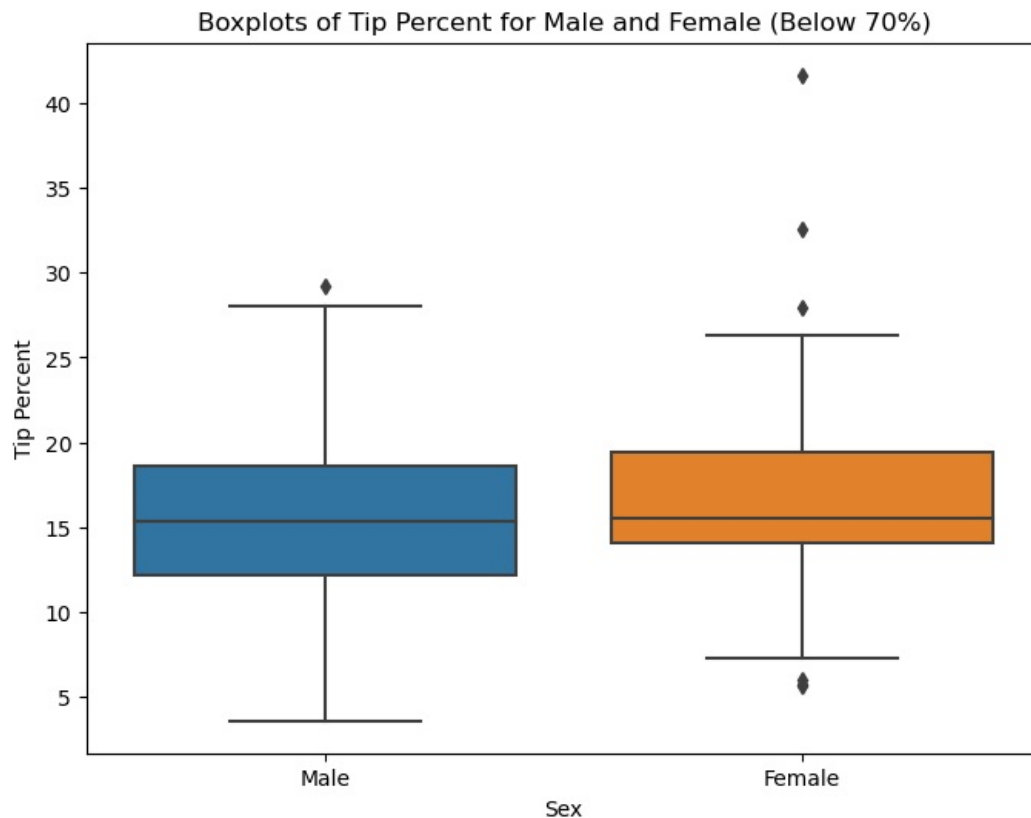
	smoker	count_lunch	count_dinner	percent_lunch	percent_dinner
0	No	45	106	66.176471	60.227273
1	Yes	23	70	33.823529	39.772727

```
In [11]: #3.4
import seaborn as sns
```

```
import matplotlib.pyplot as plt
plt.figure(figsize=(6, 4))
sns.boxplot(x='sex', y='tip', data=tips_df)
plt.title('Boxplots of Tips for Male and Female')
plt.xlabel('Sex')
plt.ylabel('Tip')
plt.show()
```



```
In [3]: #3.5
import seaborn as sns
import matplotlib.pyplot as plt
tips_df = sns.load_dataset('tips')
tips_df['tip_percent'] = (tips_df['tip'] / tips_df['total_bill']) * 100
filtered_tips_df = tips_df[tips_df['tip_percent'] < 70]
plt.figure(figsize=(8, 6))
sns.boxplot(x='sex', y='tip_percent', data=filtered_tips_df)
plt.title('Boxplots of Tip Percent for Male and Female (Below 70%)')
plt.xlabel('Sex')
plt.ylabel('Tip Percent')
plt.show()
```



```
In [12]: #4.1
import pandas as pd
avocado_series = pd.read_csv(r"C:\Users\gated\.ipython\avocado.csv")
df = pd.read_csv(r"C:\Users\gated\.ipython\avocado.csv")
print(avocado_series)
missing_values = df.isnull()
missing_values_count = missing_values.sum()
```

```

print("Count of missing values per column:")
print(missing_values_count)
numeric_columns = df.select_dtypes(include=['number']).columns
df[numeric_columns] = df[numeric_columns].fillna(df[numeric_columns].mean())
missing_values_after_imputation = df.isnull().sum()
print("Count of missing values after mean imputation:")
print(missing_values_after_imputation)
text_columns = df.select_dtypes(include=['object']).columns
placeholder_value = "Unknown" # Placeholder value for missing text or object values
df[text_columns] = df[text_columns].fillna(placeholder_value)
missing_values_after_imputation = df.isnull().sum()
print("Count of missing values after place holder imputation:")
print(missing_values_after_imputation)

```

	Date	AveragePrice	TotalVolume	Small	Large	AllSizes	\
0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	
1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	
2	2015-12-13	0.93	118220.22	794.70	109149.67	NaN	
3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	
4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	
...	
18244	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	
18245	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	
18246	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	
18247	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	
18248	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	

	TotalBags	Type	Year	Region
0	8696.87	conventional	2015.0	Albany
1	9505.56	conventional	NaN	Albany
2	8145.35	conventional	2015.0	Albany
3	5811.16	conventional	2015.0	Albany
4	6183.95	conventional	2015.0	Albany
...
18244	13498.67	organic	2018.0	WestTexNewMexico
18245	9264.84	organic	2018.0	WestTexNewMexico
18246	9394.11	organic	2018.0	WestTexNewMexico
18247	10969.54	organic	2018.0	WestTexNewMexico
18248	12014.15	organic	2018.0	WestTexNewMexico

```

[18249 rows x 10 columns]
Count of missing values per column:
Date          176
AveragePrice  184
TotalVolume   192
Small         194
Large         178
AllSizes      184
TotalBags     184
Type          204
Year          196
Region        169
dtype: int64
Count of missing values after mean imputation:
Date          176
AveragePrice   0
TotalVolume    0
Small          0
Large          0
AllSizes       0
TotalBags      0
Type           204
Year           0
Region         169
dtype: int64
Count of missing values after place holder imputation:
Date          0
AveragePrice  0
TotalVolume   0
Small         0
Large         0
AllSizes      0
TotalBags     0
Type           204
Year           0
Region         0
dtype: int64

```

```

In [14]: #4.2
import pandas as pd
avocado_series = pd.read_csv(r"C:\Users\gabed\.ipython\avocado.csv")
df = pd.read_csv(r"C:\Users\gabed\.ipython\avocado.csv")
df['Type'] = df['Type'].astype('category')
df['Year'] = df['Year'].astype('category')
df['Region'] = df['Region'].astype('category')
filtered_df = df[~df['Region'].isin(['TotalUS', 'West'])]
filtered_df = filtered_df.sort_values(by='Date')
mean_price_2016 = filtered_df[filtered_df['Year'] == 2016]['AveragePrice'].mean()
mean_price_2017 = filtered_df[filtered_df['Year'] == 2017]['AveragePrice'].mean()

```



```

if mean_price_2017 > mean_price_2016:
    print("The mean price of an avocado is higher in 2017 compared to 2016.")
elif mean_price_2017 < mean_price_2016:
    print("The mean price of an avocado is higher in 2016 compared to 2017.")
else:
    print("The mean price of an avocado is the same in 2016 and 2017.")


```

The mean price of an avocado is higher in 2017 compared to 2016.


```

In [33]: #4.3
import pandas as pd
import matplotlib.pyplot as plt
avocado_series = pd.read_csv(r"C:\Users\gabel\ipynthon\avocado.csv")
sales_by_region = avocado_series.groupby('Region')['TotalVolume'].sum()
plt.figure(figsize=(10, 10))
sales_by_region.plot(kind='barh')
plt.title('TotalVolume of Avocado Sales by Region')
plt.xlabel('TotalVolume')
plt.ylabel('Region')
plt.show()
highest_sales_region = sales_by_region.idxmax()
print("State with the highest sales of avocados by volume:", highest_sales_region)
subset_state_data = avocado_series[avocado_series['Region'] == highest_sales_region]
plt.figure(figsize=(8, 6))
plt.hist(subset_state_data['AveragePrice'], bins=20, color='blue', edgecolor='black')
plt.title('Histogram of Mean Price for {}'.format(highest_sales_region))
plt.xlabel('Mean Price')
plt.ylabel('Frequency')
plt.grid(True)
plt.show()
correlation_index = subset_state_data['AveragePrice'].corr(subset_state_data['TotalVolume'])
print("Correlation index between mean price and total volume for state {}: {:.2f}".format(highest_sales_region,

```



State with the highest sales of avocados by volume: TotalUS



Correlation index between mean price and total volume for state TotalUS: -0.80

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js